

# **ECGR 3123: Data Communications & Networking**

**Tao Han, Ph.D**

**Assistant Professor**

**Department of Electrical and Computer Engineering**

**The William States Lee College of Engineering**

**University of North Carolina-Charlotte**

**<http://webpages.uncc.edu/than3>**

# **SOCKET PROGRAMMING**

# Sockets

## ➤ What is a Socket

- Socket is the API to access transport layer functions

## ➤ How to use sockets

- Setup socket
  - Where is the remote machine (IP address, hostname)
  - What service gets the data (port)
- Send and Receive
  - Designed just like any other I/O in unix
  - send -- write
  - recv -- read
- Close the socket

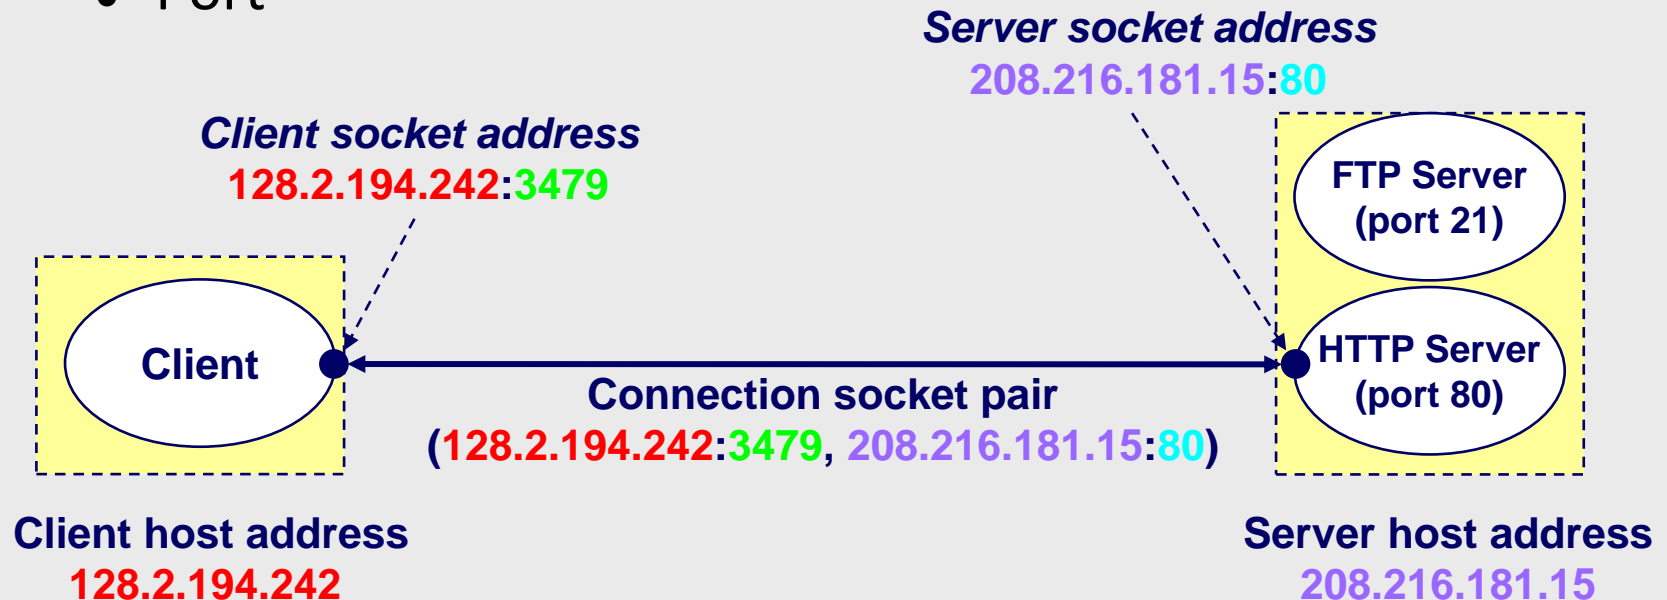
# Identify the Destination

## ➤ Addressing

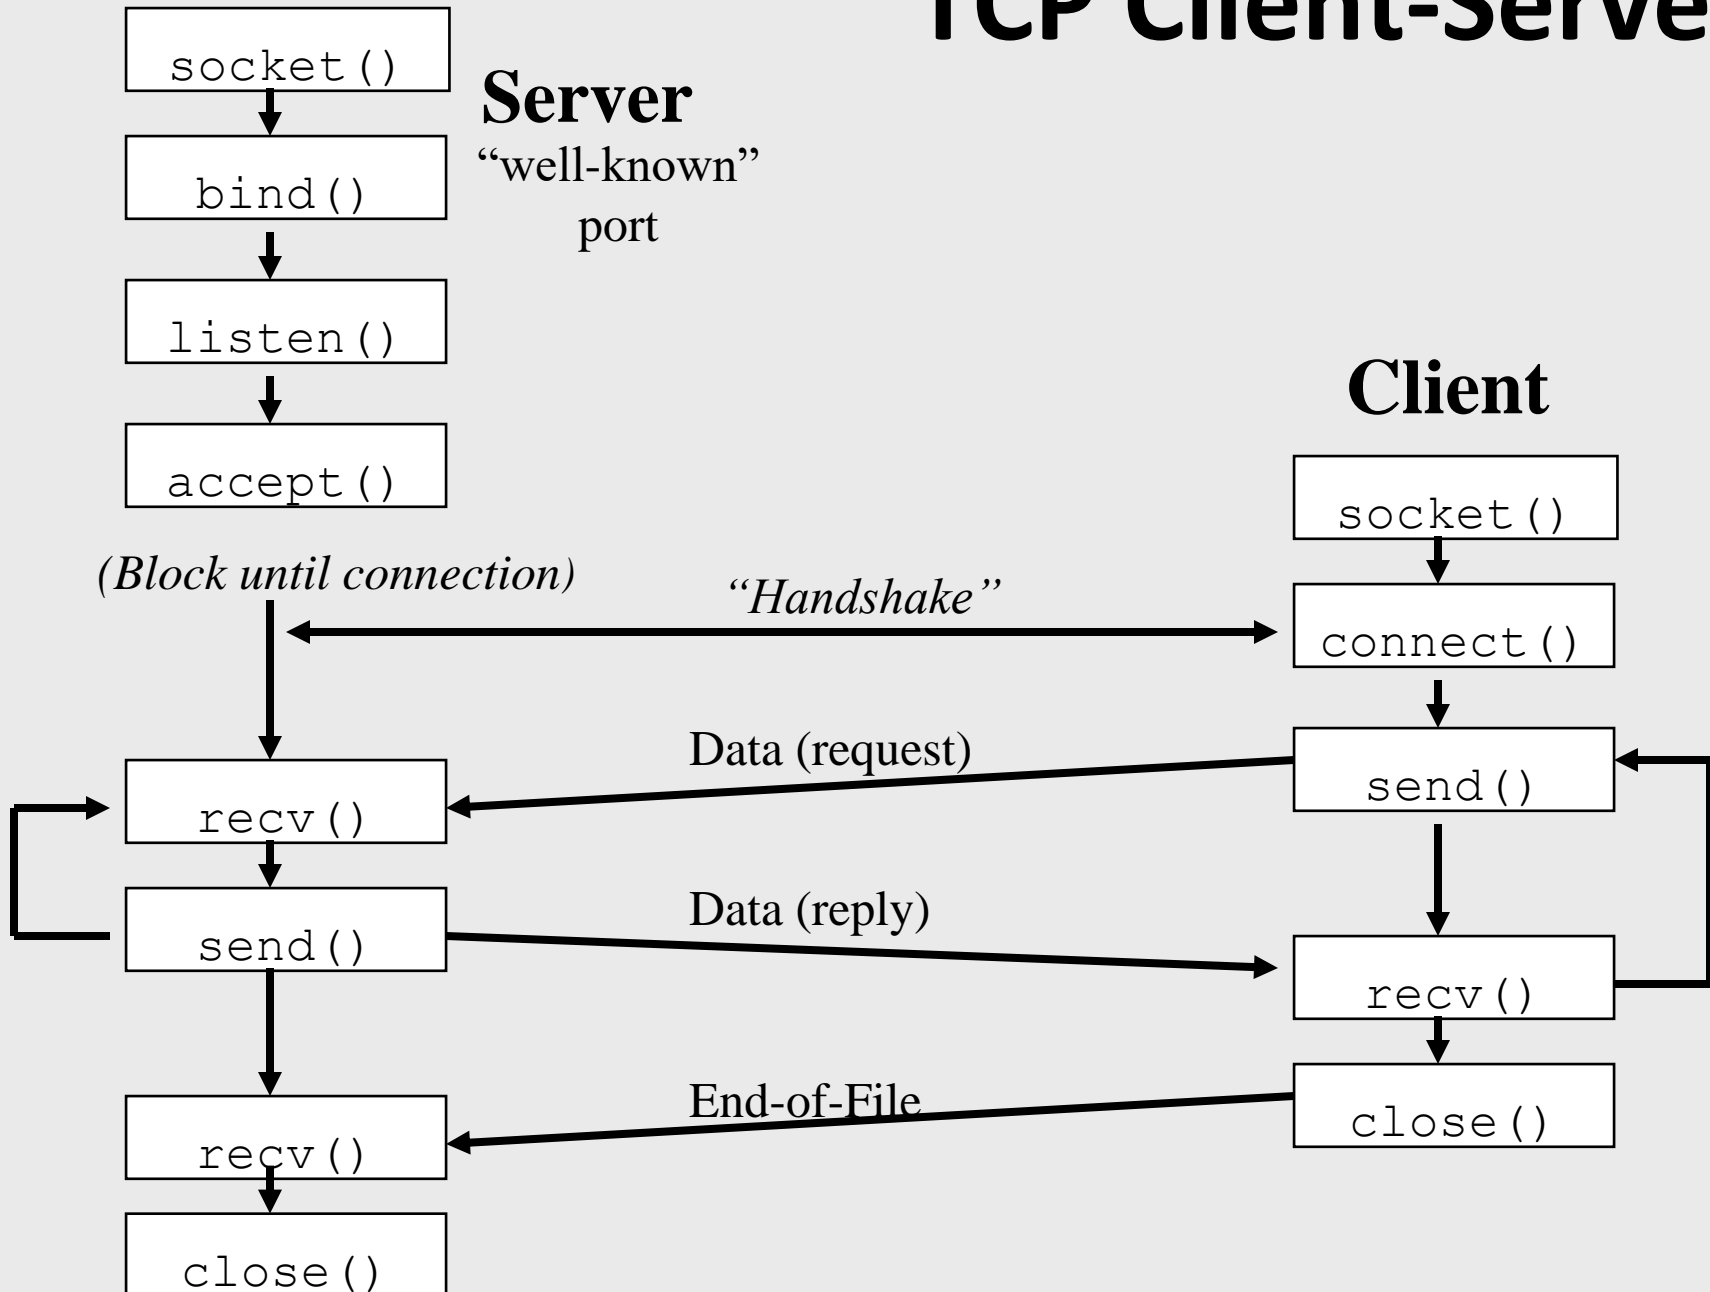
- IP address
- Hostname (resolve to IP address via DNS)

## ➤ Multiplexing

- Port



# TCP Client-Server



# Step 1 – Setup a socket()

```
int socket(int family, int type, int protocol);
```

Create a socket/file descriptor, giving access to transport layer service.

- *family* is one of
  - AF\_INET (IPv4), AF\_INET6 (IPv6), AF\_LOCAL (local Unix),
  - AF\_ROUTE (access to routing tables), AF\_KEY (new, for encryption)
- *type* is one of
  - SOCK\_STREAM (TCP), SOCK\_DGRAM (UDP)
  - SOCK\_RAW (for special IP packets, PING, etc. Must be root)
- *protocol* is 0 (used for some raw socket options)
- upon success returns socket descriptor
  - Integer, like file descriptor
  - Return -1 if failure

## Step 2: `bind()`

```
int bind(int sockfd, const struct  
sockaddr *myaddr,  
        socklen_t addrlen);
```

Assign a local protocol address (“name”) to a socket.

- *sockfd* is socket descriptor from `socket()`
- *myaddr* is a pointer to address struct with:
  - *port number* and *IP address*
  - if port is 0, then host will pick ephemeral port
  - IP address = `INADDR_ANY` (unless multiple nics)
- *addrlen* is length of structure
- returns 0 if ok, -1 on error
  - `EADDRINUSE` (“Address already in use”)

# Socket Address Structure

```
struct in_addr {  
    in_addr_t s_addr;           /* 32-bit IPv4 addresses */  
};  
  
struct sockaddr_in {  
    unit8_t      sin_len;       /* length of structure */  
    sa_family_t  sin_family;    /* AF_INET */  
    in_port_t    sin_port;      /* TCP/UDP Port num */  
    struct in_addr sin_addr;     /* IPv4 address (above) */  
    char sin_zero[8];           /* unused */  
}
```

- Are also “generic” and “IPv6” socket structures



# Convert Byte order

- **Little Endian** – In this scheme, low-order byte is stored on the starting address (A) and high-order byte is stored on the next address (A + 1).
- **Big Endian** – In this scheme, high-order byte is stored on the starting address (A) and low-order byte is stored on the next address (A + 1).

# Convert Byte order

'h' : host byte order

'n' : network byte order

's' : short (16bit)

'l' : long (32bit)

```
uint16_t htons(uint16_t);
```

```
uint16_t ntohs(uint16_t);
```

```
uint32_t htonl(uint32_t);
```

```
uint32_t ntohl(uint32_t);
```

In\_addr values should be stored in  
network byte order

# bind() example

```
int sd;  
struct sockaddr_in ma;  
sd = socket(AF_INET, SOCK_STREAM, 0);  
  
ma.sin_family = AF_INET;  
ma.sin_port = htons(5100);  
ma.sin_addr.s_addr = htonl(INADDR_ANY);  
if(bind(sd, (struct sockaddr *)&ma,  
        sizeof(ma)) != -1) { ...}
```

# Step 3: `listen()`

```
int listen(int sockfd, int backlog);
```

Change socket state for TCP server.

- *sockfd* is socket descriptor from `socket()`
- *backlog* is maximum number of *incomplete* connections
  - historically 5
  - rarely above 15 on a even moderate Web server!

# Step 4: `accept()`

```
int accept(int sockfd, struct sockaddr  
cliaddr, socklen_t *addrlen);
```

Return next completed connection.

- *sockfd* is socket descriptor from `socket()`
- *cliaddr* and *addrlen* return protocol address from client
- **returns brand new descriptor**, created by OS

# listen/connect() example

```
struct sockaddr_in ca;
```

```
//After bind..
```

```
listen(sd, 5);
```

```
calen = sizeof(ca);
```

```
cd = accept(sd, (struct sockaddr  
*) ca, &calen);
```

```
//read and write using cd
```

# Client: connect()

```
int connect(int sockfd, const struct  
sockaddr *servaddr, socklen_t addrlen);
```

Connect to server.

- *sockfd* is socket descriptor from `socket()`
- *servaddr* is a pointer to a structure with:
  - *port number* and *IP address*
  - must be specified (unlike `bind()`) – how?
- *addrlen* is length of structure
- client doesn't need `bind()`
  - OS will pick ephemeral port
- returns socket descriptor if ok, -1 on error

# connect() example

```
int sd;  
struct sockaddr_in sa;  
sd = socket(AF_INET, SOCK_STREAM, 0);  
  
sa.sin_family = AF_INET;  
sa.sin_port = htons(5100);  
sa.sin_addr.s_addr =  
    htonl(inet_addr("128.226.123.101"));  
if(connect(sd, (struct sockaddr *) &sa,  
    sizeof(sa)) != -1) { ...}
```



# close () / shutdown

```
int close(int sockfd);
```

Close socket for use.

- *sockfd* is socket descriptor from `socket ()`
- Returns -1 if error

# Sending and Receiving

```
int recv(int sockfd, void  
    *buff, size_t mbytes, int  
    flags);
```

```
int send(int sockfd, void  
    *buff, size_t mbytes, int  
    flags);
```

# Domain Name System (DNS)

- What if I want to send data to “www.slashdot.org”?
  - DNS: Conceptually, DNS is a database collection of host entries


```
struct hostent {  
    char *h_name;      // official hostname  
    char **h_aliases;  // vector of alternative hostnames  
    int  h_addrtype;   // address type, e.g. AF_INET  
    int  h_length;     // length of address in bytes, e.g. 4 for IPv4  
    char **h_addr_list; // vector of addresses  
    char *h_addr;      // first host address, synonym for h_addr_list[0]  
};
```

- hostname -> IP address
  - *struct hostent \*gethostbyname(const char \*name);*
- IP address -> hostname
  - *struct hostent \*gethostbyaddr(const char \*addr, int len, int type);*

# Project 1: Client-Server Chat

## ➤ Basic requirement (15%):

- Socket in C++ (Windows, Linux, Mac)
- Client – Server Communications: Client and server can continuously exchange messages
- One TCP socket and one UDP socket



## ➤ Additional credit (5%)

- Client – Server – Client Communications



## ➤ Additional credit (5%)

- Group chat: multiple clients can chat with each other in a group

